

# DONG JIN (JIN) YU

10905 Ohio Ave. Apt 306 ◇ Los Angeles, CA 90024 ◇ (310) · 990 · 7382

yudj0604@gmail.com ◇ github.com/ydj0604 ◇ www.linkedin.com/pub/dong-jin-jin-yu/58/939/360

## EDUCATION

---

### University of California, Los Angeles

June 2015

B.S. in Computer Science (June 2015) CS GPA: 3.80 / Overall GPA: 3.59

M.S. in Computer Science (June 2016) with focus on Distributed Systems

## PROJECT

---

### PhotoHangout (Java)

June 2015

- Real-time collaborative photo editing IOS Application (<https://www.youtube.com/watch?v=VOUcd-RqMaQ>)
- Designed/implemented back-end architecture from service layer to DB (MySQL) required for the mobile app.
- Utilized Websocket (JSR 356) for live streaming session and REST (Jersey) for back-end services.

### Content Distribution Network (C++)

January 2015

- Designed/implemented CDN with Load Balancer, Meta-data Server, Local Cache Server, Central Storage Server.
- Distributes clients' requests to cache servers based on selection algorithm to store/transfer files efficiently.
- Client syncs automatically with cloud storage, and file hash is used to optimize the sync flow.

### Bot/Human Prediction (Python, R)

May 2015

- Analyzed bid history dataset, bidder info dataset to construct an ensemble model to predict.
- Successfully extracted multiple useful features to help achieve 88% accuracy of the model.

### Hangman Solver (C++)

January 2015

- Optimized search time and memory usage by using Trie to construct a dictionary.
- Predicts a phrase correctly with less than 3 misses 85% of the time.

## WORK EXPERIENCE

---

### hulu

June 2015 - August 2015

*Software Developer Intern, Core Services*

*Los Angeles, CA*

- Designed/implemented a scalable real-time view counter for all videos/shows at hulu that handles 1000 views/s.
- Implemented it in two different ways (one with HBase and the other with Cassandra) to compare their performances.
- Understood HBase/Cassandra internal structures/algorithms and designed schemas accordingly.
- Implemented daemon that pops views from Redis-backed service, aggregates, and makes batch writes to HBase/Cassandra.
- Implemented a process that runs periodically to scan the entire HBase/Cassandra and copy all the data to MySQL.
- Used Flask to implement a simple webapp that displays view counts for demonstration, verification purpose.
- Wrote MapReduce jobs to parse logs in HDFS to analyze the performances of various existing services.
- Set up HBase stage cluster with 4 VMs from scratch and maintain it.

### Scalable Analytics Institute (UCLA)

January 2015 - June 2015

*Undergraduate Researcher*

*Los Angeles, CA*

- Implemented large-scale ML algorithm APIs (logistic regression, k-means) for IMRU (Iterative Map Reduce Update: a MapReduce framework optimized for iterative reads) on top of Hyracks (Distributed FS).
- Optimized IMRU's reducer performance by implementing an aggregation tree algorithm for reduce jobs.

### eBay Inc. (Paypal)

June 2014 - September 2014

*Software Developer Intern, Merchant Services*

*San Jose, CA*

- Developed REST service (in Java, Jersey) that updates/retrieves merchant preferences.
- Implemented DB access layer, business logic, client stub to make calls to other services.